

Marvin Abisrror Zarate

**On Selecting of  
heuristics functions for Domain–Independent  
planning.**

**Brasil**

**2015, v-1.9.5**



Marvin Abisrror Zarate

**On Selecting of  
heuristics functions for Domain–Independent planning.**

Paper presented to the Federal University of Viçosa, as part of the requirements of Graduate Computer Science program, for obtaining the title of Magister Scientiae.

Universidade de Viçosa – UFV

Centro de Ciencias Exactas e Tecnologicas (CCE)

Programa de Pós-Graduação

Supervisor: Levi Henrique Santana de Lelis

Co-supervisor: Santiago Franco

Brasil

2015, v-1.9.5

Marvin Abisrror Zarate

On Selecting of  
heuristics functions for Domain–Independent planning./ Marvin Abisrror Zarate. –  
Brasil, 2015, v-1.9.5-

73 p. : il. (algumas color.) ; 30 cm.

Supervisor: Levi Henrique Santana de Lelis

Tese (Mestrado) – Universidade de Viçosa – UFV  
Centro de Ciencias Exactas e Tecnologicas (CCE)  
Programa de Pós-Graduação, 2015, v-1.9.5.

1. Palavra-chave1. 2. Palavra-chave2. 2. Palavra-chave3. I. Orientador. II. Univer-  
sidade xxx. III. Faculdade de xxx. IV. Título

Marvin Abisrror Zarate

## **On Selecting of heuristics functions for Domain–Independent planning.**

Paper presented to the Federal University of Viçosa, as part of the requirements of Graduate Computer Science program, for obtaining the title of Magister Scientiae.

Trabalho aprovado. Brasil, 24 de novembro de 2012:

---

**Levi Henrique Santana de Lelis**  
Orientador

---

**Professor**  
Convidado 1

---

**Professor**  
Convidado 2

Brasil  
2015, v-1.9.5



*This thesis is dedicated to my Mother.*





# Acknowledgements

I would like to express my sincere gratitude to my advisor PhD. Levi Henrique Santana de Lelis, for the continuous support and guidance during the thesis process. His valuable advice, patience and encouragement have been of great importance for this work.

Besides my advisor, I would like to thank to my co—advisor: PhD. Santiago Franco for his insightful feedback, interest and tough questions.

To the professors of the DTI, particularly the Master Degree program with all its members, played an invaluable role in my graduate education.

Last, but not least, I would like to thank my Mother.



*“Não vos amoldeis às estruturas deste mundo,  
mas transformai-vos pela renovação da mente,  
a fim de distinguir qual é a vontade de Deus:  
o que é bom, o que Lhe é agradável, o que é perfeito.  
(Bíblia Sagrada, Romanos 12, 2)*



# Abstract

In this dissertation we present a greedy method based on the theory of supermodular optimization for selecting a subset of heuristics functions from a large set of heuristics with the objective of reducing the running time of the search algorithms.

([HOLTE et al., 2006](#)) showed that search can be faster if several smaller pattern databases are used instead of one large pattern database. We introduce a greedy method for selecting a subset of the most promising heuristics from a large set of heuristics functions to guide the A\* search algorithm. If the heuristics are consistent, our method selects a subset which is guaranteed to be near optimal with respect to the resulting A\* search tree size. In addition to being consistent, if all heuristics have the same evaluation time, our subset is guaranteed to be near optimal with respect to the resulting A\* running time. We implemented our method in Fast Downward and showed empirically that it produces heuristics which outperform the state of the art heuristics in the International Planning Competition benchmarks.

**Key-words:** Heuristics. selection.



# List of Figures

Figure 1	– The left tile–puzzle is the initial distribution of tiles and the right tile–puzzle is the goal distribution of tiles. Each one represent a State.	32
Figure 2	– Heuristic Search: $I$ : Initial State, $s$ : Some Sate, $G$ : Goal State . . . . .	33
Figure 3	– Out of place heuristic . . . . .	33
Figure 4	– Manhatham distance heuristic . . . . .	34
Figure 5	– One heuristic of size $M$ . . . . .	35
Figure 6	– Two heuristics of size $M/2$ . . . . .	36
Figure 7	– $N$ heuristics of size $M/N$ . . . . .	36
Figure 8	– Blocks world with three blocks. . . . .	37
Figure 9	– Sokoban with four blocks solved . . . . .	40
Figure 10	– Type system and the search space representation. . . . .	49
Figure 11	– The heuristic value is the position of the empty tile in a Specific state.	53
Figure 12	– Each row represent a different type. . . . .	54
Figure 13	– Search tree using Type System . . . . .	55





# List of Tables

Table 1	– Ratios of the number of nodes expanded using $h_{max}(\zeta')$ to the number of nodes expanded using $h_{max}(\zeta)$ . . . . .	60
Table 2	– Coverage of different planning systems on the 2011 IPC benchmarks. For the <b>RGHS</b> and <b>Max</b> approaches we also present the average number of heuristics <b>RGHS</b> selects ( $ \zeta' $ ). . . . .	64



# Contents

	<b>Introduction</b>	<b>21</b>
<b>I</b>	<b>PREPARATION OF THE RESEARCH</b>	<b>23</b>
<b>1</b>	<b>ABOUT THE PROBLEM</b>	<b>25</b>
	<i>The purpose of this section is to motivate the problem.</i>	
<b>1.1</b>	<b>Problem Statement and Motivation</b>	<b>25</b>
<b>1.2</b>	<b>Aim and Objectives</b>	<b>26</b>
1.2.1	Aim	26
1.2.2	Objectives	26
<b>1.3</b>	<b>Scope, Limitations, and Delimitations</b>	<b>26</b>
<b>1.4</b>	<b>Justification</b>	<b>27</b>
<b>1.5</b>	<b>Hypothesis</b>	<b>27</b>
<b>1.6</b>	<b>Contribution of the Thesis</b>	<b>27</b>
<b>1.7</b>	<b>Organization of the Thesis</b>	<b>27</b>
<b>II</b>	<b>LITERATURE REVIEW</b>	<b>29</b>
<b>2</b>	<b>BACKGROUND</b>	<b>31</b>
	<i>The purpose of this section is to understand the problem.</i>	
<b>2.1</b>	<b>Similar Selection Systems</b>	<b>31</b>
<b>2.2</b>	<b>Problem definition</b>	<b>31</b>
<b>2.3</b>	<b>Heuristics</b>	<b>32</b>
2.3.1	Out of place (O.P)	33
2.3.2	Manhatham Distance (M.D)	34
<b>2.4</b>	<b>Heuristic Generators</b>	<b>34</b>
2.4.1	Pattern Database (PDB)	34
<b>2.5</b>	<b>Take advantage of Heuristics</b>	<b>34</b>
<b>2.6</b>	<b>Number of heuristics created</b>	<b>35</b>
<b>2.7</b>	<b>Heuristic Subset</b>	<b>36</b>
<b>2.8</b>	<b>Problem Domains</b>	<b>37</b>
2.8.1	Blocks world	37
2.8.2	Barman	37
2.8.3	Elevators	37
2.8.4	Floortile	38



**Index** . . . . . **73**



# Introduction

This thesis is concerned with cost-optimal state-space planning using the  $A^*$  algorithm (HART P. E.; NILSSON; RAPHAEL, 1968). We assume that a pool,  $\zeta$ , of hundreds or even thousands of heuristics is available, and that the final heuristic used to guide  $A^*$ ,  $h_{max}$ , will be defined as the maximum over a subset  $\zeta'$  of those heuristics ( $h_{max}(s, \zeta') = \max_{h \in \zeta'} h(s)$ ). The choice of the subset  $\zeta'$  can hugely affect the efficiency of  $A^*$ . For a given size  $N$  and planning task  $\nabla$ , a subset containing  $N$  heuristics from  $\zeta$  is optimal if no other subset containing  $N$  heuristics from  $\zeta$  results in  $A^*$  expanding fewer nodes when solving  $\nabla$ .

Exists many problems of Artificial Intelligent (AI), such as: Finding the shortest path from one point to another in a game map, 8-tile-puzzle, Rubick's cube, etc. The level of difficulty to solve the problems mentioned are linked with the size of the search space generated.

State-space search algorithms have been used to solve the problems mentioned above. And in this dissertation we study the approach to solve problems in order to reduce the size of the search tree generated and the running time of the search algorithm using the best subset of heuristics selected from a large set of heuristics.





## Part I

### Preparation of the research



# 1 About the Problem

*The purpose of this section is to motivate the problem.*

## 1.1 Problem Statement and Motivation

Every problem of Artificial Intelligence can be cast as a state space problem. The state space is a set of states where each state represents a possible solution to the problem and each state is linked with other states if there exists a function that goes from one state to another. In the search space there are many solutions that represent the same state, each of these solutions are called nodes. So, many nodes can be represented as one state. To find the solution of the problem is required the use of search algorithms such as: Depth First Search (DFS), which looks for the solution of the problem traversing the search space exploring the nodes in each branch before backtracking up to find the solution. Another search algorithm is Breadth First Search (BFS), which looks for the solution exploring the neighbors nodes first, before moving to the next level of neighbors. The mentioned algorithms have the characteristic that when they do the search, they generate a larger search space. The search space that these algorithms generate are called Brute force search tree (BFST).

There are other types of algorithms called heuristics informed search, which are algorithms that require the use of heuristics. The heuristic is the estimation of the distance for one node in the search tree to get to the near solution. The heuristic informed search generates smaller search tree in comparison to the BFST, because the heuristic guides the search exploring the nodes that are in the solution path and prunes the nodes which are not. Also, the use of heuristics reduces the running time of the search algorithm.

There are different approaches to create heuristics, such as: Pattern Databases (PDBs), Neural Network, and Genetic Algorithm. These systems that create heuristics receive the name of Heuristics Generators. And one of the approaches that have showed most successful results in heuristic generation is the PDBs, which is memory-based heuristic functions obtained by abstracting away certain problem variables, so that the remaining problem ("pattern") is small enough to be solved optimally for every state by blind exhaustive search. The results stored in a table, represent a PDB for the original problem. The abstraction of the search space gives an admissible heuristic function, mapping states to lower bounds.

Exists many ways to take advantage of all the heuristics that can be created, for

example: (HOLTE et al., 2006) showed that search can be faster if several smaller pattern databases are used instead of one large pattern database. In addition (DOMSHLAK; KARPAS; MARKOVITCH, 2010) and (TOLPIN et al., 2013) results showed that evaluating the heuristic lazily, only when they are essential to a decision to be made in the search process is worthy in comparison to take the maximum of the set of heuristics. Then, using all the heuristics do not guarantee to solve the major number of problems in a limit time.

## 1.2 Aim and Objectives

### 1.2.1 Aim

The objective of this dissertation is to develop meta-reasoning approaches for selecting heuristics functions from a large set of heuristics with the goal of reducing the running time of the search algorithm employing these functions.

### 1.2.2 Objectives

- Demonstrate that the problem of finding the optimal subset of  $\zeta$  of size  $N$  for a given problem task is supermodular respect the size of the search tree.
- Develop an approaches to obtain the cardinality of the subsets of heuristics found.
- Develop an approach to find a subset of heuristics from a large pool of heuristics that optimize the number of nodes expanded in the process of search.
- Develop an approach for selecting a subset of heuristics functions based on the evaluation cost of each heuristic.
- Develop an strategy to drop heuristics during the sampling that do not improve the objective function.
- Use Stratified Sampling (SS) algorithm for predicting the search tree size of Iterative-Deepening A\* (IDA\*). And use SS as our utility function.

## 1.3 Scope, Limitations, and Delimitations

We implemented our method in Fast Downward (HELMERT, 2006) and the problems we want to solve are the 2011 International Planning Competition (IPC) instances.

We used the 2011 instances instead of the 2014 instances because the former do not have problems with conditional effects, which are currently not handled by PDB heuristics.

Our meta–reasoning described in this thesis are going to try to solve the major number of problems using the most promising heuristics from a large set of heuristics. The exact way to create the large set of heuristics is beyond the scope of this thesis.

## 1.4 Justification

In the last few decades, Artificial Intelligence has made significant strides in domain–independent planning. The use of heuristics search approach have contributed to problem solving, where the use of an appropriate heuristic often means substantial reduction in the time needed to solve hard problems.

That is why we propose a meta–reasoning that will try to solve the major number of problems without relying on domain knowledge, to guide the A\* search algorithm.

## 1.5 Hypothesis

This thesis will intend to prove the hypotheses listed below:

- **H1:** Prove that our objective function of selection is supermodular respect to the A\* search tree size and respect to the A\* running time.
- **H2:** Prove that SS is an effective prediction method to our meta–reasoning.

## 1.6 Contribution of the Thesis

The main contributions of this Thesis are:

- Provide a meta–reasoning approach based on the size of the search tree generated.
- Provide a meta–reasoning approach based on the evaluation cost of each heuristic.

## 1.7 Organization of the Thesis

The Thesis is organized as follows:

1. In Part I, the background of the thesis is provided which also includes our motivation and define the scope.
2. In Part II, we review the state of the art in selection.

3. In Part III, we introduce Random Greedy Heuristic Selection (RGHS) and the prediction methods.
4. In Part IV, we explain the results obtained by using GRHS and compared with other planner systems.
5. We conclude in Part V, by discussing the concluding remarks.

In the next chapter, the domain 8–tile–puzzle is used to understand the concepts that will be helpful for the other Parts.

## Part II

### Literature Review





## 2 Background

*The purpose of this section is to understand the problem.*

### 2.1 Similar Selection Systems

An optimization procedure which is similar to ours is presented by (RAYNER; STURTEVANT; BOWLING, 2013), but their procedure maximizes the average heuristic value. By contrast, the meta-reasoning we are proposing minimizes the search tree size.

Our meta-reasoning requires a prediction of the number of nodes expanded by  $A^*$  using any given subset. Although there are methods for accurately predicting the number of nodes expanded by Iterative Deepening- $A^*$  (KORF, 1985) (IDA\*). (SS system (LELIS; ZILLES; HOLTE, 2013)), these methods can't be easily adapted to  $A^*$  because  $A^*$ 's duplicate pruning makes it very difficult to predict how many nodes will occur at depth  $d$  of  $A^*$ 's search tree (the tree of nodes expanded by  $A^*$ ). As a part of our proposal, we present SS for predicting the size of the search tree.

The system most similar to ours is RIDA\* (BARLEY; FRANCO; RIDDLE, 2014). RIDA\* also selects a subset from a pool of heuristics to guide the  $A^*$  search. In RIDA\* this is done by starting with an empty subset and trying all combination of size one before trying the combination of size two and so on. RIDA\* stops after evaluating a fixed number of subsets. While RIDA\* is able to evaluate a set of heuristics with tens of elements, our meta-reasoning is able to evaluate a set of heuristics with thousands of elements.

### 2.2 Problem definition

A  $SAS^+$  planning task (BÄCKSTRÖM; NEBEL, 1995) is a 4 tuple  $\nabla = \{V, O, I, G\}$ .  $V$  is a set of *state variables*. Each variable  $v \in V$  is associated with a finite domain of possible  $D_v$ . A state is an assignment of a value to every  $v \in V$ . The set of possible states, denoted  $V$ , is therefore  $D_{v_1} \times \dots \times D_{v_2}$ .  $O$  is a set of operators, where each operator  $o \in O$  is triple  $\{pre_o, post_o, cost_o\}$  specifying the preconditions, postconditions (effects), and non-negative cost of  $o$ .  $pre_o$  and  $post_o$  are assignments of values to subsets of variables,  $V_{pre_o}$  and  $V_{post_o}$ , respectively. Operator  $o$  is applicable to state  $s$  if  $s$  and  $pre_o$  agree on the assignment of values to variables in  $V_{pre_o}$ . The effect of  $o$ , when applied to  $s$ , is to set the variables in  $V_{post_o}$  to the values specified in  $post_o$  and to set all other variables to the value they have in  $s$ .  $G$  is the goal condition, an assignment of values to a subset of variables,  $V_G$ . A state is a goal state if it and  $G$  agree on the assignment of values to the variable in  $V_G$ .  $I$  is the initial state, and the planning task,  $\nabla$ , is to find an optimal (least-cost)

sequence of operators leading from  $I$  to a goal state. We denote the optimal solution cost of  $\nabla$  as  $C^*$

The state space problem illustrated in the figure 1 is a game that consists of a frame of numbered square tiles in random order with one tile missing. The puzzle also exists in other sizes, particularly the smaller 8–puzzle. If the size is  $3 \times 3$  tiles, the puzzle is called the 8–puzzle or 9–puzzle, and if  $4 \times 4$  tiles, the puzzle is called the 15–puzzle or 16–puzzle named, respectively, for the number of tiles and the number of spaces. The object of the puzzle is to place the tiles in order by making sliding moves that use the empty space.

The legal operators are to slide any tile that is horizontally or vertically adjacent to the blank into the blank position. The problem is to rearrange the tiles from some random initial configuration into a particular desired goal configuration. The 8–puzzle contains 181,440 reachable states, the 15–puzzle contains about  $10^{13}$  reachable states, and the 24–puzzle contains almost  $10^{25}$  states.

Initial			Goal		
4	1	2	1	2	3
8		3	4	5	6
5	7	6	7	8	

Figure 1: The left tile–puzzle is the initial distribution of tiles and the right tile–puzzle is the goal distribution of tiles. Each one represent a State.

Instead of using an algorithm of Brute force search that will analyze all the possible solutions. We can obtain heuristics from the problem of the slide tile puzzle that will help us to solve the problem.

## 2.3 Heuristics

State–space algorithms, such as  $A^*$  (HART P. E.; NILSSON; RAPHAEL, 1968), are important in many AI applications.  $A^*$  uses the  $f(s) = g(s) + h(s)$  cost function to guide its search. Here,  $g(s)$  is the cost of the path from the start state  $s$ , and  $h(s)$  is the estimated cost–to–go from  $s$  to a goal;  $h(\cdot)$  is known as the heuristic function. The heuristic is the mathematical concept that represent to the estimate distance from the node  $s$  to the nearest goal state.



Figure 2: Heuristic Search:  $I$ : Initial State,  $s$ : Some Sate,  $G$ : Goal State

In the figure 2 the optimal distance from the Initial State  $I$  to the state  $s$  is 4 and represented by  $g(s)$ . The  $h^*(s)$  represent the optimal distance from  $s$  to the Goal State  $G$ . And the  $h(s)$  is the estimation distance from  $s$  to  $G$ .

A heuristic function  $h(s)$  estimates the cost of a solution path from  $s$  to a goal state. A heuristic is admissible if  $h(s) \leq h^*(s)$  for all  $s \in V$ , where  $h^*(s)$  is the optimal cost of  $s$ . A heuristic is consisten iff  $h(s) \leq c(s, t) + h(t)$  for all states  $s$  and  $t$ , where  $c(s, t)$  is the cost of the cheapest path from  $s$  to  $t$ . For example, the heuristic function provided by a pattern database (PDB) heuristic (CULBERSON; SCHAEFFER, 1998) is admissible and consistent.

Given a set of admissible and consistent heuristics  $\zeta = \{h_1, h_2, \dots, h_M\}$ , the heuristic  $h_{max}(s, \zeta) = \max_{h \in \zeta} h(s)$  is also admissible and consistent. When describing our method we assume all heuristics to be consistent. We define  $f_{max}(s, \zeta) = g(s) + h_{max}(s, \zeta)$ , where  $g(s)$  is the cost of the path expanded from  $I$  to  $s$ .  $g(s)$  is minimal when A\* using a consistent heuristic expands  $s$ . We call an A\* search tree the tree defined by the states expanded by A\* using a consistent heuristic while solving a problem  $\nabla$ .

The heuristics can be obtained from each state of the problem. For example, for the problem of the 8–tile–puzzle figure 1 we can get two heuristics.

### 2.3.1 Out of place (O.P)

Counts the number of objects out of place.




Figure 3: Out of place heuristic

The tiles numbered with 4, 1, 2, 3, 6, 7, 5, 8, and 4 are out of place then each object count as 1 and the sum would be 8.

### 2.3.2 Manhatham Distance (M.D)

Counts the minimum number of operations to get to the goal state.




Figure 4: Manhatham distance heuristic

The tile 4 count 1 to get to the goal position. The tile 1 count 1 to get to the goal position. The tile 2 count 1 to get to the goal position. The tile 3 count 1 to get to the goal position. The tile 6 count 1 to get to the goal position. The tile 7 count 1 to get to the goal position. The tile 5 count 1 to get to the goal position. The tile 8 count 1 to get to the goal position. Then the sum would be 10.

In order to solve the problem, we get the heuristics, which are information from the problem to solve the problem. Exists systems that can create heuristics for each problem. Those systems are called Heuristic Generators.

## 2.4 Heuristic Generators

Heuristic Generators works by creating abstractions of the original problem space. The approach that has showed more successful results lately is PDB.

### 2.4.1 Pattern Database (PDB)

It's obtained by abstracting away certain problem variables, so that the remaining problem ("pattern") is small enough to be solved optimally for every state by blind exhaustive search. The results stored in a table, represent a PDB for the original problem. The abstraction of the search space gives an admissible heuristic function, mapping states to lower bounds.

## 2.5 Take advantage of Heuristics

The heuristics generators can create hundreds or even thousand of heuristics. In fact, exists different ways to take advantage of those heuristics. For example: If we want to use all the heuristics created by the heuristic generator. It would not be a good idea to use all of them because the main problem involved would be the time to evaluate each

heuristic in the search tree, it could take too much time.

One way to take advantage of heuristics would be to take the maximum of the set of heuristics. For example, using three different heuristics  $h1, h2$  and  $\max(h1, h2)$ . Heuristic  $h1$  and  $h2$  are based on domain abstractions and the  $\max(h1, h2)$  is the maximum heuristic value of  $h1$  and  $h2$ .

Exists different approaches to take advantage from a large set of heuristics. In this dissertation we use the meta-reasoning based on the minimum size of the search tree generated and the minimum evaluation time.

## 2.6 Number of heuristics created

Let's suppose we have to run our meta-reasoning using  $M$  amount of memory available. The question would be: How many heuristics our system should handle in order to avoid out of memory errors? So, one of the objectives of this tesis is to find the number of heuristics that our subset  $\zeta'$  should have.

([HOLTE et al., 2006](#)) observed that maximizing over  $N$  pattern databases of size  $M/N$ , for a suitable choice  $N$ , produces a significant reduction in the number of nodes generated compared to using a single pattern database of size  $M$ .

In the Figures 6 and 7 we are taking advantage of the heuristics doing the maximization of all the heuristics created. In this thesis we are interested in heuristics that generate smaller search tree.






Figure 5: One heuristic of size  $M$

Figure 6: Two heuristics of size  $M/2$ Figure 7:  $N$  heuristics of size  $M/N$ 

## 2.7 Heuristic Subset

The heuristics generator systems can create a large number of heuristics. Let's suppose  $|\zeta| = 1000$  heuristics were created considering the time and memory available and we want to select the best  $N = 100$  heuristics. This would be:

$$\binom{1000}{100} = 10^{138} \text{possibilities}$$

So, try to select heuristics from a large set of heuristics are going to be treated as an optimization problem. Then, in order to obtain a good selection of subset of heuristics, our objective function should guarantee two properties: Monotonicity and Submodularity, that would be explained in the next Part.

## 2.8 Problem Domains

Some of the problems we are trying to solve are the optimal domains for International Planning Competition (IPC).

### 2.8.1 Blocks world

This domain consists of a set of blocks, a table and a robot hand. The blocks can be on top of other blocks or on the table; a block that has nothing on it is clear; and the robot hand can hold one block or be empty. The goal is to find a plan to move from one configuration of blocks to another.




Figure 8: Blocks world with three blocks.

The solution shown in Figure 8 is to unblock number 1 from block number 2; stack block number 2 on block number 1; finally, stack block number 3 on block number 2. We are interested in optimal or near-optimal solutions for this kind of problem.

Depending on the number  $n$  of blocks this domain can have very large state spaces. For example, using  $n = 20$  this domain has approximately  $10^{20}$  different states.

### 2.8.2 Barman

In this domain there is a robot barman that manipulates drink dispensers, glasses and a shaker. The goal is to find a plan of the robot's actions that serves a desired set of drinks.

### 2.8.3 Elevators

The idea for this domain came up from the Miconic domain of IPC2, however the domain has been designed from scratch. The scenario is the following: There is a building

with  $N+1$  floors, numbered from 0 to  $N$ . The building can be separated in blocks of size  $M+1$ , where  $M$  divides  $N$ . Adjacent blocks have a common floor. For example, suppose  $N=12$  and  $M=4$ , then we have 13 floors in total (ranging from 0 to 12), which form 3 blocks of 5 floors each, being 0 to 4, 4 to 8 and 8 to 12.

The building has  $K$  fast (accelerating) elevators that stop only in floors that are multiple of  $M/2$  (so  $M$  has to be an even number). Each fast elevator has a capacity of  $X$  persons. Furthermore, within each block, there are  $L$  slow elevators, that stop at every floor of the block. Each slow elevator has a capacity of  $Y$  persons (usually  $Y < X$ ).

There are costs associated with each elevator starting/stopping and moving. In particular, fast (accelerating) elevators have negligible cost of starting/stopping but have significant cost while moving. On the other hand, slow (constant speed) elevators have significant cost when starting/stopping and negligible cost while moving. Travelling times between floors are given for any type of elevator, taking into account the constant speed of the slow elevators and the constant acceleration of the fast elevators.

There are several passengers, for which their current location (i.e. the floor they are) and their destination are given. The planning problem is to find a plan that moves the passengers to their destinations while it maximizes some criterion.

#### 2.8.4 Floortile

A set of robots use different colors to paint patterns in floor tiles. The robots can move around the floor tiles in four directions (up, down, left and right). Robots paint with one color at a time, but can change their spray guns to any available color. However, robots can only paint the tile that is in front (up) and behind (down) them, and once a tile has been painted no robot can stand on it.

For the IPC set, robots need to paint a grid with black and white, where the cell color is alternated always. This particular configuration makes the domain hard because robots should only paint tiles in front of them, since painting tiles behind make the search to reach a dead-end.

#### 2.8.5 Nomystery

In this domain, a truck moves in a weighted graph; a set of packages must be transported between nodes; actions move along edges, and load/unload packages; each



move consumes the edge weight in fuel. In brief, Nomystery is a straightforward problem similar to the ones contained in many IPC benchmarks.

### 2.8.6 Openstacks

The openstacks domain is based on the "minimum maximum simultaneous open stacks" combinatorial optimization problem, which can be stated as follows: A manufacturer has a number of orders, each for a combination of different products, and can only make one product at a time.

The total required quantity of each product is made at the same time (because changing from making one product to making another requires a production stop). From the time that the first product included in an order is made to the time that all products included in the order have been made, the order is said to be "open" and during this time it requires a "stack" (a temporary storage space). The problem is to order the making of the different products so that the maximum number of stacks that are in use simultaneously, or equivalently the number of orders that are in simultaneous production, is minimized (because each stack takes up space in the production area).

### 2.8.7 ParcPrinter

This domain models the operation of the multi-engine printer, for which one prototype is developed at the Palo Alto Research Center (PARC). This type of printer can handle multiple print jobs simultaneously. Multiple sheets, belonging to the same job or different jobs, can be printed simultaneously using multiple Image Marking Engines (IME). Each IME can either be color, which can print both color and black and white images, or mono, which can only print black and white image. Each sheet needs to go through multiple printer components such as feeder, transporter, IME, inverter, finisher and need to arrive at the finisher in order.

### 2.8.8 Parking

This domain involves parking cars on a street with  $N$  curb locations, and where cars can be double-parked but not triple-parked. The goal is to find a plan to move from one configuration of parked cars to another configuration, by driving cars from one curb location to another. The problems in the competition contain  $2^{*(N-1)}$  cars, which allows one free curb space and guarantees solvability.

### 2.8.9 Sokoban

This domain is inspired by the popular Sokoban puzzle game where an agent has the goal of pushing a set of boxes into specified goal locations in a grid with walls. The competition problems are generated in a way that guarantees solvability and generally are easy problem instances for humans.




Figure 9: Sokoban with four blocks solved

In the next Part, we will introduce the meta-reasoning proposed for selecting heuristics and will expand on the properties of our objective functions.

## Part III

### Approach Proposal



### 3 Meta-Reasoning for selection

*The purpose of this section is to introduce the meta-reasoning proposed.*

#### 3.1 Random Greedy Heuristic Selection (RGHS)

We present a random greedy algorithm selection for approximately solving the heuristic subset selection problem while optimizing different objective functions. We consider the following general optimization problem.

$$\begin{aligned} & \text{minimize}_{\zeta' \in 2^{|\zeta|}} \Psi(\zeta', \nabla) \\ & \text{subject to } |\zeta'| = N \end{aligned} \tag{3.1}$$

Where  $\Psi(\zeta', \nabla)$  is an objective function and  $N$  is the desired subset size.  $N$  could be determined by a hard constraint such as the maximum number of PDBs one can store in memory. According to (RAYNER; STURTEVANT; BOWLING, 2013) it is unlikely that there is an efficient algorithm for solving Equation 3.1. We use an algorithm based on the work of (BUCHBINDER et al., 2014) we call Random Greedy Heuristic Selection (RGHS) to approximately solve Equation 3.1 for different functions  $\Psi$ .

---

**Algoritmo 1:** Random Greedy Heuristic Selection

---

**Input** : problem  $\nabla$ , set of heuristics  $\zeta$ , cardinality  $N$

**Output** : heuristic subset  $\zeta' \subseteq \zeta$  of size  $N$

```

1  $\zeta'_0 \leftarrow \emptyset$ 
2 for  $i = 1$  to  $N$  do
3   Let  $M_i \subseteq \zeta \setminus \zeta'_{i-1}$  be a subset of size  $N$  minimizing  $\sum_{h \in M_i} \Psi(\zeta'_{i-1} \cup \{h\}) - \Psi(\zeta'_i)$ 
4   Let  $h_i$  be a uniformly random element from  $M_i$ 
5    $\zeta'_i \leftarrow \zeta'_{i-1} \cup \{h_i\}$ 
6 return  $\zeta'_N$ 
```

---

Algorithm 1 shows RGHS. RGHS receives as input a problem  $\nabla$ , a set of heuristics  $\zeta$ , a cardinality size  $N$ , and it returns a subset  $\zeta' \subseteq \zeta$ . In each iteration  $i$  RGHS randomly selects a heuristic  $h_i$  from a pool  $M_i$  of "good" heuristics and adds  $h_i$  to  $\zeta'$ .  $M_i$  is defined as follows.  $M_i \subseteq \zeta \setminus \zeta'_{i-1}$  is a subset of size  $N$  minimizing  $\sum_{h \in M_i} \Psi(\zeta'_{i-1} \cup \{h\}) - \Psi(\zeta'_i)$ , where  $\zeta'_{i-1}$  is the subset of heuristics RGHS selects prior to iteration  $i$  of the algorithm.  $M_i$  contains the  $N$  heuristics that when individually combined with  $\zeta'_{i-1}$  minimizes  $\Psi$  the most. RGHS returns  $\zeta'$  once it reaches the desired size  $N$ .

The work of (RAYNER; STURTEVANT; BOWLING, 2013) uses a greedy algorithm introduced by (NEMHAUSER; WOLSEY; FISHER, 1978) for approximately solving the heuristic subset selection problem. In contrast with the RGHS presented above, which is based on the work of (BUCHBINDER et al., 2014), (NEMHAUSER; WOLSEY; FISHER, 1978)’s approach does not offer near-optimality guarantees when the problem’s objective functions is non-monotone. As mentioned before  $A^*$ ’s running time is a non-monotone objective function. While RGHS also offers guarantees for non-monotone objective functions, it retains the same guarantees offered by (NEMHAUSER; WOLSEY; FISHER, 1978)’s algorithm when optimizing monotone objective functions (BUCHBINDER et al., 2014). That is why we only consider (BUCHBINDER et al., 2014)’s approach in our theoretical and experimental analyses.

### 3.2 Approximately Minimizing Search Tree Size

The first objective function  $\Psi$  we consider accounts for the number of expansions  $A^*$  performs while solving a given planning problem. When solving  $\nabla$  using the consistent heuristic function  $h_{max}(\zeta')$  for  $\zeta' \subseteq \zeta$ ,  $A^*$  expands in the worse case  $J(\zeta', \nabla)$  nodes, where

$$J(\zeta', \nabla) = |\{s \in V \mid f_{max}(s, \zeta') \leq C^*\}| \quad (3.2)$$

$$J(\zeta', \nabla) = |\{s \in V \mid h_{max}(s, \zeta') \leq C^* - g(s)\}| \quad (3.3)$$

We write  $J(\zeta')$  or simply  $J$  instead of  $J(\zeta', \nabla)$ . RGHS is guaranteed to find near-optimal solutions when we use  $J$  as the objective function  $\Psi$ , as we now demonstrate. In the following analysis all heuristic functions are assumed to be consistent. We also assume that  $A^*$  expands all nodes  $n$  with  $f(n) \leq C^*$  while solving  $\nabla$ , as shown in Equation 3.2.

**Lemma 3.2.1.**  $J(\zeta' \cup \{h\}) \leq J(\zeta')$  for any  $\zeta'$  and any  $h$ .

*Proof.* Fix  $\zeta'$  and  $h$ . Then

$$\begin{aligned} J(\zeta' \cup \{h\}) &= |\{s \in V \mid h_{max}(s, \zeta' \cup \{h\}) \leq C^* - g(s)\}| \\ &\leq |\{s \in V \mid h_{max}(s, \zeta') \leq C^* - g(s)\}| \\ &= J(\zeta') \end{aligned}$$

Where the inequality follows from the fact that  $h_{max}(s, \zeta' \cup \{h\}) \geq h_{max}(s, \zeta')$  for all  $s$ .

Let  $S$  be a set and  $\phi$  a function over  $2^S$ .  $\phi$  is supermodular if for any  $A, B, x$  with  $A \subseteq B \subseteq S$  and  $x \in S \setminus B$ :

$$\phi(A) - \phi(A \cup \{x\}) \geq \phi(B) - \phi(B \cup \{x\}) \quad (3.4)$$

Intuitively, Equation 3.4 captures the idea of diminishing returns. In the context of search tree size, if we add a heuristic function  $h$  to a set of heuristics  $A$  strictly contained in a set  $B$ , then we would expect, then we would expect  $J(A) - J(A \cup \{h\})$  to be larger than  $J(B) - J(B \cup \{h\})$  as  $h$  would "contribute more" to  $A$  than to  $B$ .

**Lemma 3.2.2.**  $J$  is supermodular.

*Proof.* Let  $A \subset B \subset \zeta$  and  $h \in \zeta \setminus B$ . By Lemma 3.2.1,  $J(A) - J(A \cup \{h\}) \geq 0$  and  $J(B) - J(B \cup \{h\}) \geq 0$ . We consider two cases.

Case 1.  $J(B) - J(B \cup \{h\}) = 0$ . Then  $J(A) - J(A \cup \{h\}) \geq 0$  yields  $J(A) - J(A \cup \{h\}) \geq J(B) - J(B \cup \{h\})$ .

Case 2.  $J(B) - J(B \cup \{h\}) = k > 0$ . Let  $s_1, \dots, s_k$  be all the states in  $V$  that satisfy  $h_{\max}(s_i, B) \leq C^* - g(s_i)$  and  $h_{\max}(s_i, B \cup \{h\}) > C^* - g(s_i)$ . This implies  $h(s_i) > C^* - g(s_i)$ , for all  $i \in \{1, \dots, k\}$ . Further, since  $A \subset B$ , we have  $h_{\max}(s_i, A) \leq h_{\max}(s_i, B) \leq C^* - g(s_i)$ , for all  $i \in \{1, \dots, k\}$ . Consequently,  $J(A) - J(A \cup \{h\}) \geq k = J(B) - J(B \cup \{h\})$ .

Lemma 3.2.1 and 3.2.2 are sufficient for using a result by (BUCHBINDER et al., 2014) to conclude the following.

**Theorem 3.2.3.** Let  $\zeta'$  be a subset selected by GHS. Then  $J(\zeta', \nabla)$  is within a factor of  $\frac{e+1}{e} \approx 1.36$  of optimal.

### 3.3 Approximately Minimizing A\*'s Running Time

Another objective function  $\Psi$  we consider accounts for the A\* running time and is defined as follows. Let  $T(\zeta', \nabla)$  be an approximation to the running time of A\* when using  $h_{\max}(\zeta')$  for solving  $\nabla$ , defined as follows.

$$T(\zeta', \nabla) = J(\zeta', \nabla) \cdot t_{h_{\max}}(\zeta') \quad (3.5)$$

where, for any heuristic function  $h$ , the term  $t_h$  refers to the running time used for computing the  $h$ -value of any state  $s$ .

We assume that  $t_h$  to be independent of  $s$ , which is a reasonable assumption for several heuristics such as PDBs.

In order to compute the running time of  $A^*$  exactly we would also have to account for the time required for node generation and for the operations on  $A^*$ 's **OPEN** and **CLOSED** lists. However, these two factors do not depend directly on the heuristic employed. Thus,  $T(\zeta', \nabla)$  is reasonable approximation for  $A^*$ 's running time for the heuristic subset selection problem.

**Theorem 3.3.1.** *Suppose  $t_{h_i} = t_{h_j}$  for any  $h_i$  and  $h_j \in \zeta$ . Then for any fixed subset size  $N$ , **RGHS** yields a subset  $\zeta'$  that is within a factor  $\frac{e+1}{e}$  of optimal with respect to  $T(\zeta', \nabla)$*

*Proof.* Since  $t_h$  is constant over  $h \in \zeta$ , the value  $t_{h_{max}}(\zeta')$  is independent of  $\zeta'$ . Hence the value  $T(\zeta', \nabla)$  is a constant factor of  $J(\zeta', \nabla)$ . The latter is within a factor of  $\frac{e+1}{e}$  of optimal by Theorem 3.2.3.

The assumption that  $t_{h_i} = t_{h_j}$  for any  $h_i, h_j \in \zeta$  often does not hold in domain-independent planning. For example, the **iPDB** heuristic (HASLUM et al., 2007) can be few order of magnitude faster than the Incremental **LM-Cut** (HELMERT; DOMSH-LAK, 2009) heuristic. In order to lift such an assumption we first define  $A^*$ 's running time in terms of its computational cost as follows.

$$T(\zeta', \nabla) = J(\zeta', \nabla) + \beta(\zeta') \quad (3.6)$$

Here  $\beta(\zeta')$  is the computational cost incurred by using  $h_{max}(\zeta')$ .  $T(\zeta', \nabla)$  accounts for the computational cost of expanding  $J(\zeta', \nabla)$  nodes added of the computational cost of employing a set of heuristics  $\zeta'$ .

We assume that the computational cost  $\beta(\zeta' \cup \{h\}) = \beta(\zeta') + \beta(h)$ , i.e., the computational cost of employing heuristic  $h$  during search is constant and independent of other heuristics in  $\zeta'$ . Although this assumption is unlikely to hold in practice as  $\beta(h)$  depends on the number of times  $A^*$  uses  $h$  to evaluate nodes during search, which in turn depends on the heuristics in  $\zeta' \setminus h$ , we expect that the differences of  $\beta(h)$ -values to be negligible for different  $\zeta'$  sets.

$T'$  is clearly non-monotone as the reduction of  $J'(\zeta', \nabla)$  caused by the addition of a heuristic to  $\zeta'$  might not compensate for the increase in  $\beta(\zeta')$ . We now show that  $T'$  is supermodular.

**Lemma 3.3.2.**  *$T$  is supermodular.*

*Proof.* Let  $A \subset B \subseteq \zeta$  and  $h \in \zeta \setminus B$ . We need to show that  $T'(A) - T'(A \cup \{x\}) \geq T'(B) - T'(B \cup \{x\})$ . We have that  $T'(A) - T'(A \cup \{x\}) = J(A) - J(A \cup \{x\}) - \beta(x)$ . and



$T'(B) - T'(B \cup \{x\}) = J(B) - J(B \cup \{x\}) - \beta(x)$ . Thus, we have that  $J(A) - J(A \cup \{x\}) \geq J(B) - J(B \cup \{x\})$ , which according to Lemma 3.2.2 is supermodular.

Since  $T'$  is supermodular, another result by (BUCHBINDER et al., 2014) and Lemma 3.3.2 allow us to conclude the following.

**Theorem 3.3.3.** *Let  $\zeta'$  be a subset selected by RGHS. Then  $T(\zeta', \nabla)$  is within a factor of  $\frac{e+1}{e} \approx 1.63$  of optimal.*

### 3.4 Estimating Tree Size and Running Time

In practice RGHS used approximations of  $J$ ,  $T$ , and  $T'$  instead of their exact values. This is because computing  $J$ ,  $T$ , and  $T'$  exactly would require solving  $\nabla$ . We denote the approximations of  $J$  as  $\hat{J}$ , and since both  $T$  and  $T'$  model  $A^*$ 's running time, we denote the approximation for both as  $\hat{T}$ .

We use the Culprit Sampler (CS) introduced by (BARLEY; FRANCO; RIDDLE, 2014) and the Stratified Sampling (SS) algorithm introduced by (CHEN, 1992) for computing  $\hat{J}$  and  $\hat{T}$ . Each of the two algorithms has its strengths and weaknesses, which we explore in the experimental Part.

Both CS and SS must be able to quickly estimate the values of  $\hat{J}(\zeta')$  and  $\hat{T}(\zeta')$  for any subset  $\zeta'$  of  $\zeta$  so they can be used in RGHS's optimization process.

### 3.5 Culprit Sampler (CS)

CS runs a time-bounded  $A^*$  search while sampling  $f$ -culprits and  $b$ -culprits to estimate the values of  $\hat{J}$  and  $\hat{T}$ .

**Definition 3.5.1.** (*f-culprit*) Let  $\zeta = \{h_1, h_2, \dots, h_M\}$  be a set of heuristics. The  $f$ -culprit of a node  $n$  in an  $A^*$  search tree is defined as the tuple  $F(n) = \langle f_1(n), f_2(n), \dots, f_M(n) \rangle$ , where  $f_i(n) = g(n) + h_i(n)$ . For any  $n$ -tuple  $F$ , the counter  $C_F$  denotes the number of nodes  $n$  in the tree with  $F(n) = F$ .

**Definition 3.5.2.** (*b-culprit*) Let  $\zeta = \{h_1, h_2, \dots, h_M\}$  be a set of heuristics and  $b$  a lower bound on the solution cost  $\nabla$ . The  $b$ -culprit of a node  $n$  in an  $A^*$  search tree is defined as the tuple  $B(n) = \langle y_1(n), y_2(n), \dots, y_M(n) \rangle$ , where  $y_i(n) = 1$  if  $g(n) + h_i(n) \leq b$  and  $y_i(n) = 0$ , otherwise. For any binary  $n$ -tuple  $B$ , the counter  $C_B$  denotes the number of nodes  $n$  in the tree with  $B(n) = B$ .

**CS** works by running an  $A^*$  search bounded by a user-specified time limit. Then, **CS** compresses the information obtained in the  $A^*$  search (i.e., the  $f$ -values of all nodes expanded according to all heuristics  $h$  in  $\zeta$ ) in b-culprits, which are later used for computing  $\hat{J}$ . The bf-culprits are generated as an intermediate step for computing the b-culprits, as we explain below. The maximum number of f-culprits and b-culprits in an  $A^*$  search tree is equal to the number of nodes in the tree expanded by the time-bounded  $A^*$  search. However, in practice the number of f-culprits is usually much lower than the number of nodes in the tree. Moreover, in practice, the total number of different b-culprits tends to be even lower than the total number of f-culprits. Given a planning problem  $\nabla$  and a set of heuristics  $\zeta$ , **CS** samples the  $A^*$  search tree as follows.

- 1.- **CS** runs  $A^*$  using  $h_{min}(s, \zeta) = \min_{h \in \zeta} h(s)$  until reaching a user-specified time limit.  $A^*$  using  $h_{min}$  expands node  $n$  if it were to expand  $n$  while using any of the heuristics in  $\zeta$  individually. For each node  $n$  expanded in this time-bounded search we store  $n$ 's f-culprit and its counter.
- 2.- Let  $f_{maxmin}$  be the largest  $f$ -value according to  $h_{min}$  encountered in the time-bounded  $A^*$  search described above. We now compute the set  $\mathbb{B}$  of b-culprits and their counters based on the f-culprits and on the value of  $f_{maxmin}$ . This is done by iterating over all f-culprits once.

The process described above is performed only once **RGHS**'s execution. The value of  $\hat{J}(\zeta', \nabla)$  for any subset  $\zeta'$  of  $\zeta$  is then computed by iterating over all b-culprits  $\mathbf{B}$  and summing up the relevant values of  $C_B$ . The relevant values of  $C_B$  represent the number of nodes  $A^*$  would expand in a search bounded by  $b$  if using  $h_{max}(\zeta')$ . This computation can be written as follows.

$$\hat{J}(\zeta', \nabla) = \sum_{\mathbf{B} \in \mathbb{B}} W(\mathbf{B}) \quad (3.7)$$

Where  $W(\mathbf{B})$  is 0 if there is a heuristic in  $\zeta'$  whose  $y$ -value in  $\mathbf{B}$  is zero (i.e., there is a heuristic in  $\zeta'$  that prunes all nodes compressed into  $\mathbf{B}$ ), and  $C_B$  otherwise. If the time-bounded  $A^*$  search with  $h_{min}$  expands all nodes  $n$  with  $f(n) \leq C^*$ , then  $\hat{J} = J$ . In practice, however, our estimate  $\hat{J}$  will tend to be much lower than  $J$ .

The value of  $\hat{T}$  is computed by multiplying  $\hat{J}$  by the sum of the evaluation time of each heuristic in  $\zeta'$ . The evaluation time of the heuristics in  $\zeta'$  is measured in a separate process, before executing **CS**, by sampling a small number of nodes from  $\nabla$ 's start state.

## 3.6 Stratified Sampling (SS)

Stratified Sampling is a prediction algorithm that estimate the number of nodes expanded by some heuristic.

(KNUTH, 1975) created a method to estimate the size of the search tree such as IDA\*. It works doing random walk from the root of the tree. Knuth's assumption is that all branches have the same structure. So, performing a random walk down one branch is enough to estimate the size of the search tree. However, the method does not work well for unbalanced search tree. (CHEN, 1992) solved this problem with a stratification of the search tree through a *type system* to reduce the variance of the sampling process (LELIS; ZILLES; HOLTE, 2013). In the figure 10 each state of the search space is mapped to the *Type System*




Figure 10: Type system and the search space representation.

### 3.6.1 Type System

The *Type System* is a partition of the states in the state space. It is calculated based of any property of each node in the search tree. (LELIS, 2013)

A common misconception is think of *type system* as state—space abstractions. (PRIEDITIS, 1993) defines a state—space abstraction as a simplified version of the problem

in which:

- The cost of the least-cost path between two abstracted states must be less than or equal to the cost of the least-cost path between the corresponding two states in the original state-space.
- Goal states in the original state-space must be goal states in the abstracted state-space.

In contrast with state-space abstractions, a *type system* does not have these two requirements. A *type system* is just a partition of the nodes in the search tree.

The *type system* can not be represented as a graph since *type system* does not necessarily define relation between the types.

The relation between *type system* and abstractions is the following: The *type system* can not necessarily be used as abstractions, abstractions can always be used as *type system*.

### 3.6.2 Search Space and Search Tree

(LELIS; ZILLES; HOLTE, 2013) defines the search space and search tree in the following way: Let the *underlying search tree* (*UST*) be the full brute-force tree created from a connected, undirected and implicitly defined *underlying search graph* (*USG*) describing a state space. Some search algorithms expand a subtree of the *UST* while searching for a solution (e.g., a portion of the *UST* might not be expanded due to heuristic guidance); we call this subtree the *expanded search tree* (*EST*).

Let  $G = (N, E)$  be a graph representing an *ESG* where  $N$  is its set of states and for each  $n \in N$ ,  $Op(n) = \{op_i | (n, n_i) \in E\}$  is its set of operators.

**Definition 3.6.1.** *Type System* Let  $S = (N, E)$  be a *UST*.  $T = \{t_1, \dots, t_n\}$  is a *type system* for  $S$  if it is a disjoint partitioning of  $N$ . If  $n \in N$  and  $t \in T$  with  $n \in t$ , we write  $T(n) = t$ .

**SS** is a general method for approximating any function of the form  $\varphi = \sum_{n \in S} z(n)$ , where  $z$  is any function assigning a numerical value to a node.  $\varphi$  represents a numerical property of the search tree rooted at  $n^*$ . For instance, if  $z(n) = 1$  for all  $n \in S$ , then  $\varphi$  is the size of the tree.

Instead of traversing the entire tree and summing all  $z$ -values, **SS** assumes subtrees rooted at nodes of the same type will have equal values of  $\varphi$  and so only one node of each type, chosen randomly, is expanded. This is the key to **SS**'s efficiency since the search trees

of practical interest have far too many nodes to be examined exhaustively.

Given a search tree  $S$  and a type system  $T$ , SS estimates  $\varphi$  as follows. First, it samples the tree and returns a set  $A$  of *representative – weight* pairs, with one such pair for every unique type seen during sampling. In the pair  $\langle s, w \rangle$  in  $A$  for type  $t \in T$ ,  $n$  is the unique node of type  $t$  that was expanded during search and  $w$  is an estimate of the number of nodes type  $t$  in the tree.  $\varphi$  is then approximated by  $\hat{\varphi}$ , defined as,  $\hat{\varphi} = \sum_{\langle s, w \rangle \in A} w \times z(n)$ .

By making  $z(n) = 1$  for all  $n \in S$  SS produces an estimate  $\hat{J}$  of  $J$ . Similarly to our approach with CS, we obtain  $\hat{T}$  by multiplying  $\hat{J}$  by the heuristic evaluation time.

---

**Algorithm 2:** SS, a single probe

---

**Input** : root  $n^*$  of a tree and a type system  $T$   
**Output** : an array of sets  $A$ , where  $A[i]$  is the set of pairs  $\langle n, w \rangle$  for the nodes  $n$  expanded at level  $i$ .

```

1  $A[0] \leftarrow \{\langle n^*, 1 \rangle\}$ 
2  $i \leftarrow 0$ 
3 while  $A[i]$  is not empty do
4   for each element  $\langle n, w \rangle$  in  $A[i]$  do
5     for each child  $\hat{n}$  of  $n$  do
6       if  $g(\hat{n}) + h(\hat{n}) \leq d$  then
7         if  $A[i+1]$  contains an element  $\langle n', w' \rangle$  with  $T(n') = T(\hat{n})$  then
8            $w' \leftarrow w' + w$ 
9           with probability  $w/w'$ , replace  $\langle n', w' \rangle$  in
10           $A[i+1]$  by  $\langle \hat{n}, w' \rangle$ 
11         else
12           insert new element  $\langle \hat{n}, w' \rangle$  in  $A[i+1]$ 
13    $i \leftarrow i + 1$ 

```

---

In SS the types are required to be partially ordered: a node's type must be strictly greater than the type of its parent. This can be guaranteed by adding the depth of a node to the type system and then sorting the types lexicographically. That is why in our implementation of SS types at one level are treated separately from types at another level by the division of  $A$  into groups  $A[i]$ , where  $A[i]$  is the set of representative–weight pairs for the types encountered at level  $i$ . If the same type occurs on different levels the occurrences will be treated as if they were different types – the depth of search is implicitly included into all of our type systems.

Algorithm 2 shows SS in detail. Representative nodes from  $A[i]$  are expanded to

get representative nodes for  $A[i + 1]$  as follows.  $A[0]$  is initialized to contain only the root of the search tree to be probed, with weight 1 (Line 1). In each iteration (Lines 4 through 11), all nodes in  $A[i]$  are expanded. The children of each node in  $A[i]$  are considered for inclusion in  $A[i + 1]$  if their  $f$ -value do not exceed an upper bound  $d$  provided as input to **SS**. If a child  $\hat{n}$  has a type  $t$  that is already represented in  $A[i + 1]$  by another node  $n'$ , then a *merge* action on  $\hat{n}$  and  $n'$  is performed. In a merge action we increase the weight in the corresponding representative–weight pair of type  $t$  by the weight  $w(n)$  of  $\hat{n}$ 's parent  $n$  (from level  $i$ ) since there were  $w(n)$  nodes at level  $i$  that are assumed to have children of type  $t$  at level  $i + 1$ .  $\hat{n}$  will replace  $n'$  according to the probability shown in Line 9. (CHEN, 1992) proved that this probability reduces the variance of the estimation. Once all the states in  $A[i]$  are expanded, we move to the next iteration.

One run of the **SS** algorithm is called a *probe*. (CHEN, 1992) proved that the expected value of  $\hat{\varphi}$  converges to  $\varphi$  in the limit as the number of probes goes to infinity. As (LELIS; STERN; STURTEVANT, 2014), **SS** is not able to detect duplicated nodes in its sampling process. As a result, since  $A^*$  does not expanded duplicates, **SS** usually overestimates the actual number of nodes  $A^*$  expands. Thus, in the limit, as the number of probes grows large, **SS**'s prediction converges to a number which is likely to overestimate the  $A^*$  search tree size. We test empirically whether **SS** is able to allow **RGHS** to make good subset selects despite being unable to detect duplicated nodes during sampling.

Similarly to **CS**, we also define a time–limit to run **SS**. We use **SS** with an iterative–deepening approach in order to ensure an estimate of  $\hat{J}$  and  $\hat{T}$  before reaching the time limit. We set the upper bound  $d$  to the heuristic value of the start state and, after performing  $p$  probes, if there is still time, we increase  $d$  to twice its previous value. The values of  $\hat{J}$  and  $\hat{T}$  is given by the prediction produced for the last  $d$ -value in which **SS** was able to perform all  $p$  probes.

**SS** must also be able to estimate the values of  $\hat{J}(\zeta')$  and  $\hat{T}(\zeta')$  for any subset  $\zeta'$  of  $\zeta$ . This is achieved by using **SS** to estimate b–culprits (See Definition 3.5.2) instead of the search tree size directly. Similarly to **CS**, **SS** used  $h_{min}$  of the heuristics in  $\zeta$  to decide when to prune a node (See Line 6 2) while sampling. This ensures that **SS** expands a node  $n$  if  $A^*$  employing at least one of the heuristics in  $\zeta$  would expand  $n$  according to bound  $d$ . The  $C_B$  counter of each b–culprit  $B$  encountered during **SS**'s probe is given by,

$$C_B = \sum_{\langle n, w \rangle \in A \wedge B(n) = B} w \quad (3.8)$$

We recall that to compute  $B(n)$  for node  $n$  one needs to define a bound  $b$ . Here we

use the bound  $d$  used by **SS**. The average value of  $C_B$  across  $p$  probes is used to predict the search tree size for a given subset  $\zeta'$ . As explained for **CS**, this can be done by traversing over all b-culprits once.

### 3.7 8-tile-puzzle Case Using type system

In this thesis we use *type system* based only in heuristics. (ZAHAVI et al., 2010) use the simplest heuristic-based *type system* in which two nodes  $n$  and  $n'$  are of the same type if they have the same heuristic value.

C	E	C
E	M	E
C	E	C

1	2	3
4	5	6
7	8	

Figure 11: The heuristic value is the position of the empty tile in a Specific state.

Let's explain how *type system* works through an example. The problem of 8-tile-puzzle in the Figure 11, the center tile is labeled with the letter **M**, the corners labeled with **C** and the mediums with **E**. For this problem, we can define the *type system* based on the position of the empty tile regarding the position of the empty tile in the goal state. For this case, two nodes  $n$  and  $n'$  would be of the same type if  $n$  and  $n'$  have the empty tile with the same letter and with the same distance to the empty tile of the goal state.

In the Figure 12, each row represent a type. The first board shows the empty tile in the center with distance to the empty tile in the goal state equal to 2. The shortest path to get to the empty tile in the goal state would be doing: down-left or right-down. Which both moves represent the same cost. Then, the type would be (2,M). In the next board, the empty tile is in the left top, and use the letter **C**. The minimum distance is 4 because to get the goal empty tile is necessary to do: down-down-right-right or right-right-down-down, etc. So the type would be (4, C). In the next board, there are two tiles that have the same type, because both have the same letter **M** with distance equal to 3. Then, both have the type (3,E). The fourth row have two boards with the same type, because both have the same letter **C** with distance equal to 2. Then, both have the type (2, C). The fifth row have two boards with the same type, because both have the same letter **E** with distance equal to 1. Then, both have the type (1, E). In the last row the empty tile

is in the goal empty tile. So, the distance is zero and the letter is **C**. The type would be (0, **C**).




Figure 12: Each row represent a different type.



### 3.8 SS step by step

In the Figure 13, we can see how *Type System* works. In the Level 1, we have the root node, the  $w$  initialized with one. Let's suppose that three nodes are generated by the root node in the Level 2. The nodes in the Level 2 have the following types: red, blue and red from left to right respectively, and each node receive the same  $w$  of the father. In the Level 2 we apply SS's assumption, two nodes in the same level that have the same type (The same color) root subtrees of the same size and so only one node of each type must be chosen randomly to be expanded. There are two nodes with type red in Level 2. In that way, we choose randomly one of them. Let's suppose we choose the right red node. Then, we have to update the number of nodes with the type red using the  $w$ , both red node types have  $w = 1$ , then we sum the  $w$  and the new  $w = 2$ . As a result, in the Level 2 we will have two nodes of red type and one node with blue type.




Figure 13: Search tree using Type System

When nodes in the Level 2 are expanded. The blue node expands one node of type blue and the red node expands two nodes of type red and blue. The question here is how many nodes would be generated in the Level 3? The answer is:  $1 \times \text{blue} + 2 \times \text{red} + 2 \times \text{blue}$ . So, in the Level 3 we will have 2 nodes of red type and 3 nodes of type blue.

In the Level 3 the  $w$  of the node blue would have the same  $w$  of the father. The father has  $w = 1$ , then the child has  $w = 1$ . The  $w$  of the red type and blue type would be 2. Once the  $w$  has been updated for each node in the Level 3 we apply the SS's assumption again. There are two nodes with type blue. So, we choose randomly one of them and update their  $w$ . Let's choose the right blue type and the updated  $w$  would be 3 because 1 from the left blue type plus the 2 from the right blue type.

When nodes in the Level 3 are expanded. They are expanded in the following

way: The red node expands two nodes of types red and blue and the blue node expands one of red type. How many nodes would be generated at Level 4, then? The answer is:  $2 \times \text{red} + 3 \times \text{red} + 2 \times \text{blue}$ . Therefore, in the Level 4, we will have five nodes of type red and two nodes of type blue.

Finally, the number of nodes expanded in the search tree is obtained summing all  $w$  plus one (The root node). As a result, the number of nodes expanded in the search tree would be  $15 + 1 = 16$ .

### 3.9 Comparison between SS and IDA

SS is an algorithm that estimates the number of nodes expanded performed by heuristic search algorithm seeking solutions in state space. We apply SS to predict the number of nodes expanded by IDA\* in a given  $f$ -layer when using a consistent heuristics.

We first ran IDA\* for Fast-Downward benchmark for optimal domains. Our evaluation metric is coverage, i.e., number of problems solved within 24 hours time limit. We note that in 24 hours non all the instances for a specific domain using a consistent heuristic can be solved. Afterwards, run SS using as a threshold the  $f$ -layer for each instance of each domain, this process is executed using different number of probes i.e., 1, 10, 100, 1000 and 5000.

$$\frac{\sum_{s \in PI} \frac{Pred(s,d) - R(s,d)}{R(s,d)}}{|PI|} \quad (3.9)$$

Where  $PI$  is the set of problem instances,  $Pred(s,d)$  and  $R(s,d)$  are the predicted and actual number of nodes expanded by IDA\* for start state  $s$  and cost bound  $d$ . A perfect score according to the measure is 0.00.

## Part IV

# Empirical Evaluation



## 4 Empirical Evaluation

*The purpose of this section is to present how our meta-reasoning is applied when solving IPC domain problems.*

RGHS is guaranteed to find a near-optimal heuristic subset to guide the A\* search granted that it is able to compute the objective function of interest. Thus, the practical effectiveness of RGHS depends on its ability of finding good approximations  $\hat{J}$  and  $\hat{T}$ . Moreover, RGHS's effectiveness depends on the value of  $N$  and on the quality of the set of heuristics  $\zeta$ . In order to verify its practical effectiveness, we have implemented RGHS in Fast Downward (HELMERT, 2006) and tested the A\* performance using subsets of heuristics selected by RGHS while minimizing different objective functions.

We run two sets of experiments. In the first set we verify whether the approximations  $\hat{J}$  and  $\hat{T}$  provided by CS and SS allow RGHS to make near-optimal subset selections. In the second set of experiments we test the effectiveness of RGHS by measuring the total number of problem instances solved by A\* using a heuristic subset selected by RGHS.

In our experiments we implemented an approximation procedure to automatically choose the subset size selected by RGHS, which we now explain. Initially we fix  $N$  to a value (in our experiments we use  $N = 25$ ). Then, in every iteration, when computing  $M_i$ , RGHS removes from  $\zeta$  the heuristics that cannot help reducing  $\varphi$ . That is, in iteration  $i$ , all heuristics in  $\zeta \setminus \zeta'_{i-1}$  that cannot reduce  $\varphi$  when combined with  $\zeta'_{i-1}$  are removed from  $\zeta$ . RGHS stops if  $\zeta$  is empty and returns the current subset  $\zeta'$  with  $|\zeta'| < N$ . Likewise, should heuristic  $h$  improve the objective function being optimized, then we also allow RGHS to add  $h$  to  $\zeta'$  even if  $|\zeta'| > N$ . We observed in preliminary experiments that RGHS is robust to the initial value of  $N$ , as the total number of problems solved by A\* while guided by a heuristic subset selected by RGHS varied little with  $N$ .

In all our experiments we use a type-system that assigned the same type for a node with the same  $f$ -value. Such a type system has shown to be effective in guiding SS to produce accurate tree size predictions in other application domains (LELIS; ZILLES; HOLTE, 2013; LELIS; OTTEN; DECHTER, 2014).

We ran our experiments on the 2011 International Planning Competition (IPC) instances. We used the 2011 instances instead of the 2014 instances because the former do not have problems with conditional effects, which are currently not handled by PDB heuristics. All experiments are run on 2.67 GHz machines with 4GB, and are limited to

Table 1: Ratios of the number of nodes expanded using  $h_{max}(\zeta')$  to the number of nodes expanded using  $h_{max}(\zeta)$

Domain	SS		CS		$ \zeta $	n
	Ratio	$ \zeta' $	Ratio	$ \zeta' $		
Barman	2.15	27.26	2.34	30.68	4790.16	20
Elevators	1.09	18.00	1.04	20.00	155.00	1
Floortile	1.01	45.86	1.01	41.14	136.57	13
Openstacks	1.00	1.00	1.00	1.00	316.62	12
Parking	1.01	7.05	1.01	7.37	17.53	18
Pegsol	1.00	32.00	1.00	48.33	78.00	2
Scanalyzer	1.01	37.29	1.07	21.43	52.14	6
Tidybot	1.03	4.06	1.00	8.19	3259.00	15
Transport	7.33	9.70	1.45	12.90	156.50	9
Visitall	1.00	95.00	1.02	48.00	218.00	2
Woodworking	2.39	83.00	611.72	32.00	346.00	5

1,800 seconds of running time.

## 4.1 Empirical Evaluation of $\hat{J}$ and $\hat{T}$

We test whether the approximation  $\hat{J}$  provided by CS and SS allows RGHS to make near-optimal subset selections. This test is made by comparing  $J(\zeta')$  with  $J(\zeta)$ , which is minimal. The condition  $J(\zeta') \leq \alpha \cdot J(\zeta)$ , is sufficient to show that  $J(\zeta')$  is within  $\alpha$  times optimal with respect to all subsets of any size, for some constant  $\alpha$ . We are particularly interested in observing if  $J(\zeta')$  is within 1.36 of  $J(\zeta)$ .

In contrast with objective function  $J$ , there is no easy way to find the minimum of  $T$  for a subset of fixed size in general. We experiment then with the special case in which all heuristics in  $\zeta$  have the same evaluation time. This way we are able to test whether the estimates  $\hat{T}$  are allowing RGHS to make near-optimal subset selections while minimizing the  $A^*$  running time. This is because by only selecting heuristics which have the same evaluation time, if RGHS is making near-optimal subset selections with respect to  $J$ , then RGHS must also be making near-optimal subset selections with respect to  $T$  for a fixed subset size (See Theorem 3.3.1).

We collect values of  $J(\zeta)$  and  $J(\zeta')$  as follows. For each problem instance  $\nabla$  in our test set we generate a set of PDB heuristics using the GA-PDB algorithm (EDELKAMP, 2007) as described by (BARLEY; FRANCO; RIDDLE, 2014) — we call each PDB generated by this method a GA-PDB. We chose to use GA-PDBs in this experiment because they all have nearly the same evaluation time and will allow us to verify whether RGHS is making

near-optimal selections not only when minimizing  $J$  but also when minimizing  $T$ , as explained above. The number of **GA-PDBs** generated is limited in this experiment by 1,200 seconds and 1GB of memory. Also, all **GA-PDBs** we generate have 2 millions entries each. The **GA-PDBs** generated form out  $\zeta$  set. **RGHS** then selects a subset  $\zeta'$  of  $\zeta$ . Finally, we use  $h_{max}(\zeta')$  and  $h_{max}(\zeta)$  to independently try to solve  $\nabla$ . We call the system which uses  $A^*$  with  $h_{max}(\zeta)$  the **Max** approach. For **RGHS** we allow 600 seconds for selecting  $\zeta'$  and for running  $A^*$  with  $h_{max}(\zeta')$ , and for **Max** we allow 600 seconds for running  $A^*$  with  $h_{max}(\zeta)$ . Since we used 1,200 seconds to generate the heuristics, both **Max** and **RGHS** were allowed 1,800 seconds in total for solving each problem. In this experiment we test both **CS** and **SS**.

In this experiment we refer to the approach that runs  $A^*$  guided by a heuristic subset selected by **RGHS** using **CS** as **RGHS+CS**. Similarly, we write **RGHS+SS** when **SS** is used as predictor to make the heuristic subset selection.

Table 1 shows the average ratios of  $J(\zeta')$  to  $J(\zeta)$  for both **SS** and **CS** in different problem domains. The value of  $J$ , for a given problem instance, is computed as the number of nodes expanded up to the largest  $f$ -layer which is fully expanded by all approaches tested (**Max**, **RGHS** using **SS** and **RGHS** using **CS**). We only present results for instances that are not solved during **RGHS**'s **CS** sampling process. The column "n" shows the number of instances used to compute the averages of each row. We also show the average number of **GA-PDBs** generated ( $|\zeta|$ ) and the average number of **GA-PDBs** selected by **RGHS** ( $|\zeta'|$ ). This experiment shows that for most of the problems **RGHS**, using either **CS** or **SS**, is selecting a near-optimal subset of  $\zeta$ . For example, in Tidybot **RGHS** selects only a few **GA-PDBs** out of thousands when using either **SS** or **CS**. Moreover, the resulting  $A^*$  search tree size is on average at most 3% larger than optimal for **RGHS+SS**, and is optimal for **RGHS+CS**.

The exceptions in Table 1 are the ratios for Barman, Transport and Woodworking. In Transport **SS** has an average ratio of 7.33 and **CS** of 1.45. By looking at the ratios of **SS** for individual instances of **SS** for individual instances of Transport (results now show in Table 1), we noticed that **SS** is able to make optimal selections for all but one of the 9 instances considered in this experiment. Since we do not know a priori what is the instance's optimal solution cost, **SS** samples nodes with  $f$ -values much larger than the instance's optimal solution cost. We believe that, in this particular instance of Transport, by sampling a portion of the state space that is not expanded during the actual  $A^*$ , **SS** is biasing the subset selection to select heuristics that do not contribute to reducing the actual  $A^*$  search tree size.

The **SS**'s ability of sampling deep into the search space is not always harmful. For

example, **SS** allows **RGHS** to make good selections for instances of the Woodworking domain. By contrast, **CS**'s systematic approach to sampling only allow a shallow sample of the  $A^*$  search tree. As a result, **RGHS** makes a limited selection of heuristics to guide  $A^*$  search. While **RGHS** using **SS** selects an average of 83 heuristics in Woodworking instances, **RGHS** using **CS** selects only an average of 32 heuristics. This difference on sampling strategies reflects on the number of problems solved by  $A^*$ . While **RGHS+SS** solves 15 instances of the Woodworking domain, **RGHS+CS** solves only 11. In total, out of the 280 instances of the IPC 2011 benchmark set, **RGHS+SS** solves 200 problem instances in this experiment, while **RGHS+CS** only solves 193 problem instances. (The numbers of instances solved are not shown in Table 1).

## 4.2 Comparison with Other Planning Systems

The objective of this second set of experiments is to test the quality of the subset of heuristics **RGHS** selects while optimizing different objective functions. Our evaluation metric is coverage, I.E., number of problems solved within a 1,800 second time limit. We note that the 1,800-second limit includes the time to generate  $\zeta$ , select  $\zeta'$ , and run  $A^*$  using  $h_{max}(\zeta')$ . The  $\zeta$  set of heuristics is composed of a number of different **GA-PDBs**, a **PDB** heuristic produced by the **iPDB** method (HASLUM et al., 2007) and the **LM-Cut** heuristic. The generation of **GA-PDBs** is limited by 600 seconds and 1GB of memory. We use one fourth of 600 seconds to generate **GA-PDBs** with each of the following number of entries:  $\{2 \cdot 10^3, 2 \cdot 10^4, 2 \cdot 10^5, 2 \cdot 10^6\}$ . Our approach allows one to generate up to thousands of **GA-PDBs**. For every problem instance, we use exactly the same  $\zeta$  set for **Max** and all **RGHS** approaches.

## 4.3 Systems Tested

**RGHS** is tested while minimizing the  $A^*$  search tree size (**Size**) and the  $A^*$  running time (**Time**). We also use **RGHS** to maximize the sum of heuristic values in the state space (**Sum**), as suggested by (RAYNER; STURTEVANT; BOWLING, 2013). (RAYNER; STURTEVANT; BOWLING, 2013) assumed that one could uniformly sample states in the state space in order to estimate the sum of the heuristic values for a given heuristic subset. Since we are not aware of any method to uniformly sample the state space of domain-independent problems, we adapted the (RAYNER; STURTEVANT; BOWLING, 2013)'s method by using **SS** to estimate the sum of heuristic values in the search tree rooted at  $\nabla$ 's start state. We write **Size + SS** to refer to the approach that used  $A^*$  guided by a heuristic selected by **RGHS** while minimizing an estimate of the search tree size provided by **SS**. We follow the same pattern to name the other possible combinations of objective functions



and prediction algorithms (E.G., **Time+CS**).

In addition to experimenting with all combination of prediction algorithms (**CS** and **SS**) and objective functions (**Time**, **Size**), we also experiment with an approach that minimizes both the search tree size and the running time as follows. First we create a pool of heuristics  $\zeta$  composed solely of **GA-PDB** heuristics, then we apply **RGHS** while minimizing tree size and using **SS** as predictor. As explained above, in this setting **RGHS** minimizes  $J$  and  $T$  simultaneously, as all heuristics in  $\zeta$  have the same evaluation time. We call the selection of a subset of **GA-PDBs** as the *first selection*. Once the first selection is made, we test all possible combinations of the resulting  $h_{max}(\zeta')$  added to the **iPDB** and **LM-Cut** heuristics while minimizing the running time as estimated by **CS**—we call this step the *second selection*. We call the overall approach **Hybrid**.

The intuition behind **Hibrid** is that we apply **RGHS** with its strongest settings, I.E., according to Theorem 3.3.1 **RGHS** makes selections that are within the 1.36 factor of optimal with respect to both  $J$  and  $T$  when selecting from a pool of heuristics with the same evaluation time. After such a selection is made, we reduce the size of the pool of heuristics from possible thousands to only tree (the maximum of a subset of the initial **GA-PDBs**, **iPDB**, and **LM-Cut**). With only three heuristics we are able to choose the exact combination that minimizes the  $A^*$  running time the most. The reason we chose to use **SS** instead of **CS** for the first selection in **Hybrid** is that the former is able to make better subset selections in this setting, as suggested by the results discussed in the previous Part 3. Finally, as we show below, **CS** is more effective if one is interested in minimizing the  $A^*$  running time while selecting from a pool of heuristic with different evaluation times. That is why we use **CS** as predictor for the second selection in **Hybrid**.

We compare the coverage of the **RGHS** approaches with several other state-of-the-art planners. Namely, we experiment with **RIDA\*** (BARLEY; FRANCO; RIDDLE, 2014), two variants of **StoneSoup** (**StSp1** and **StSp2**) as described by (NISSIM; HOFFMANN; HELMERT, 2011), two versions of **Symba** (**SY1** and **SY2**), and  $A^*$  being independently guided by the maximum of all heuristics in  $\zeta$  (**Max**), **iPDB**, **LM-cut** and **Merge & Shrink(M&S)** (NISSIM; HOFFMANN; HELMERT, 2011).

The results are presented in Table 2. The results for the **RGHS** approaches are averages computed over 10 independent runs of the planner; the average numbers are truncated to two decimal places in our table. The variance of the results is small, thus we omitted them from the table of results.

## 4.4 Discussion of the Results

The system that solves the largest number of instances is **Hybrid**— it solves 218.4 problems on average. As explained above, we combine in **Hybrid** the strengths of both **SS** and **CS** in a single system. **SS** is used to greedily select heuristics from a pool of heuristics with similar evaluation time, and only then **CS** is used for selecting heuristics with different evaluation times. This strategy has proven particularly effective on the Barman domain where **Hybrid**’s first selection is able to select good subsets of **GA-PDBs** and its second selection is able to recognize that it must not include the **iPDB** and **LM-Cut** heuristics to the subset selected by its first selection. As a result, **Hybrid** solves more problems on this domain than any other **RGHS** approach.

**Time+CS** also performed well in our experiments—the approach solves 215 problems on average. Clearly **Hybrid** and **Time+CS** are far superior to all other approaches tested. For example, **Size + SS** and **Sum** solves only 204 and 203 problems, respectively. While minimizing the search tree size or maximizing the sum of heuristic values, **RGHS** will tend to add accurate heuristics to the selected subset, independently of their evaluation time. As a result, if not minimizing the running time, **RGHS** often adds the **LM-Cut** heuristic to  $\zeta'$  as **LM-Cut** is often the heuristic that is able to reduce the most the search tree size and to increase the most the sum of heuristic values. However, **LM-Cut** is very computationally expensive, and in various cases the search is faster if **LM-Cut** is not in  $\zeta'$ . Both **Hybrid** and **Time+CS** are able to recognize when **LM-Cut** should not be included in  $\zeta'$  because they account for the heuristics’ evaluation time.

Table 2: Coverage of different planning systems on the 2011 IPC benchmarks. For the **RGHS** and **Max** approaches we also present the average number of heuristics **RGHS** selects ( $|\zeta'|$ ).

Domains	Hybrid	CS		SS		Sum	RIDA*	SY1	SY2	StSp1	StSp2	Max	iPDB	LM-Cut	M&S
		Time	Size	Time	Size										
Barman	6.9	4.75	4	4	4	4	4	10	11	4	4	4	4	4	4
Elevators	19	19	19	19	19	19	19	20	20	18	18	19	17	18	12
Floortile	14	14	14	14	14	14	14	14	14	14	14	14	8	14	10
Nomystery	20	20	20	20	20	19	20	16	16	20	20	20	19	14	18
Openstacks	17	17	15	17	15	15	15	20	20	17	17	11	17	15	17
Parcprinter	17.7	18	19	15	16	16	18	17	17	18	18	18	16	17	16
Parking	7	7	1	2	2	2	7	2	1	5	5	2	7	2	7
Pegsol	19	19	19	19	19	19	19	19	20	19	19	19	20	17	19
Scanalyzer	14	13.87	13	12	13	13	14	9	9	14	14	14	10	12	11
Sokoban	20	20	20	20	20	20	20	20	20	20	20	20	20	20	20
Tidybot	15.9	16	16	16	16	16	17	15	17	16	16	15	14	16	9
Transport	14	13.37	11	13	13	13	10	10	11	7	8	9	8	6	7
Visitall	18	18	17	17	17	17	18	12	12	16	16	18	16	10	16
Woodworking	15.7	15	15	12	16	16	15	20	20	15	15	16	9	15	9
Total	218.4	215	203	200	204	203	210	204	208	203	204	199	185	180	175

Note that the difference on the number of problems solved by **Time+CS** and **Time+SS**: While the former solves 215 instances, the latter solved only 200. We conjecture that this happens because **SS** is not able to detect duplicated nodes during sampling. As a result, **SS** often overestimates by several orders of magnitude the actual  $A^*$ 's running time. Similarly to the **Size** and **Sum** approaches, due to **SS**'s overestimations, **Time+SS** often mistakenly adds the accurate but expensive **LM-Cut** heuristic in cases where the  $A^*$  search would be faster without **LM-Cut**'s guidance. For example, although **iPDB** tends to prune fewer nodes than **LM-Cut** in Parking instances, **iPDB** is the heuristic of choice in that domain. This is because its evaluation time is much smaller than **LM-Cut**'s. **Time+CS** solves 7 Parking instances on average as it correctly selects **iPDB** and leaves **LM-Cut** out of  $\zeta'$ . By contrast, likely due to its prediction overestimation, **Time+SS** wrongly estimates that **LM-Cut** will reduce overall search time and adds the heuristic to its selected subset. Notice, that **Size+CS** and **Size+SS** also does poorly Parking instances as they also always select **LM-Cut**.

**RIDA\*** is the most similar system to **RGHS**, as it also selects a subset of heuristics from a pool of heuristics by using an evaluation method similar to **CS**. **RIDA\*** uses a systematic approach for selecting a subset of heuristics. Namely, it starts with an empty subset and evaluates all subsets of size  $i$  before evaluating subsets of size  $i + 1$ . This procedure allows **RIDA\*** to consider only tens of heuristics in their pool. By contrast, **RGHS** is able to consider thousands of heuristics while making its selection.

The ability to handle large set of heuristics can be helpful, even if most of the heuristics in the set are redundant with each other.—as is the case with the **GA-PDBs**. The process of generating **GA-PDBs** is stochastic, thus one increases the chances of generating helpful heuristic by generating a large number of them. **RGHS** is an effective method for selecting a small set of informative heuristics from a large set of mostly uninformative ones. This is illustrated in Table 2 on the Transport domain. Compared to systems which use multiple heuristics (**StSp1** and 2, and **RIDA\***), **Time+CS** solves the largest number of Transport instances, which is due to the selection of a few key **GA-PDBs**.

The best **RGHS** approach, **Hybrid**, substantially outperforms the number of instances solved by **Max**—**Hybrid** solves on average more than 19 instances than **Max**. Finally, **Hybrid** and **Time+CS** substantially outperforms all other approaches tested, with **RIDA\*** being the closest competitor with 210 instances solved.



Part V

Conclusion



## 5 Concluding Remarks

This dissertation showed that the problem of finding the optimal subset of size  $N$  of a set of heuristics  $\zeta$  for a given problem task is supermodular with respect to the  $A^*$  search tree size and, under mild assumptions, with respect to the  $A^*$  running time. Thus, the **RGHS** algorithm which selects heuristics from  $\zeta$  one at a time is guaranteed to produce a subset  $\zeta'$  such that the number of nodes expanded by  $A^*$  while guided by the heuristic  $\zeta'$  is no more than approximately 1.36 times optimal. Furthermore, if all heuristics in  $\zeta$  have the same evaluation time, then we have the same near-optimal guarantee with respect to running time. If the heuristics in  $\zeta$  have different evaluation times, then the resulting  $A^*$  running time is no more than approximately 1.63 times optimal. In addition to minimizing the search tree size and the running time, we also experimented with an objective function that accounts for the sum of heuristic values in the state-space, as suggested by (RAYNER; STURTEVANT; BOWLING, 2013).

Since we cannot compute the values of the objective functions exactly, **RGHS** effectiveness depends on the quality of the approximations we can obtain. We tested two prediction algorithms, **CS** and **SS**, for estimating the values of the objective functions and showed empirically that both **CS** and **SS** allow **RGHS** to make near-optimal subset selections with respect to the search tree size and running time.

Finally, experiments on optimal domain-independent problems showed that **RGHS** minimizing approximations of the  $A^*$  running time outperformed all the other approaches tested, which demonstrates the effectiveness of our method for the heuristic subset selection problem.





# Bibliography

- BÄCKSTRÖM, C.; NEBEL, B. Complexity results for sas+ planning. *Computational Intelligence*, Wiley Online Library, v. 11, n. 4, p. 625–655, 1995. Citado na página [31](#).
- BARLEY, M. W.; FRANCO, S.; RIDDLE, P. J. Overcoming the utility problem in heuristic generation: Why time matters. *Proceedings of the Twenty-Fourth International Conference on Automated Planning and Scheduling, ICAPS 2014, Portsmouth, New Hampshire, USA, June 21-26, 2014*, 2014. Citado 4 vezes nas páginas [31](#), [47](#), [60](#), and [63](#).
- BUCHBINDER, N. et al. Submodular maximization with cardinality constraints. In: SIAM. *Proceedings of the Twenty-Fifth Annual ACM-SIAM Symposium on Discrete Algorithms*. [S.l.], 2014. p. 1433–1452. Citado 4 vezes nas páginas [43](#), [44](#), [45](#), and [47](#).
- CHEN, P.-C. Heuristic sampling: A method for predicting the performance of tree searching programs. *SIAM Journal on Computing*, p. 295–315, 1992. Citado 3 vezes nas páginas [47](#), [49](#), and [52](#).
- CULBERSON, J. C.; SCHAEFFER, J. Pattern databases. *Computational Intelligence*, Wiley Online Library, v. 14, n. 3, p. 318–334, 1998. Citado na página [33](#).
- DOMSHLAK, C.; KARPAS, E.; MARKOVITCH, S. To max or not to max: Online learning for speeding up optimal planning. In: *AAAI*. [S.l.: s.n.], 2010. Citado na página [26](#).
- EDELKAMP, S. Automated creation of pattern database search heuristics. In: *Model Checking and Artificial Intelligence*. [S.l.]: Springer Berlin Heidelberg, 2007. p. 35–50. Citado na página [60](#).
- HART P. E.; NILSSON, N. J.; RAPHAEL, B. A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics SSC*, IEEE, v. 4, n. 2, p. 100–107, 1968. Citado 2 vezes nas páginas [21](#) and [32](#).
- HASLUM, P. et al. Domain-independent construction of pattern database heuristics for cost-optimal planning. *AAAI*, v. 7, p. 1007–1012, 2007. Citado 2 vezes nas páginas [46](#) and [62](#).
- HELMERT, M. The fast downward planning system. *J. Artif. Intell. Res.(JAIR)*, v. 26, p. 191–246, 2006. Citado 2 vezes nas páginas [26](#) and [59](#).
- HELMERT, M.; DOMSHLAK, C. Landmarks, critical paths and abstractions: what’s the difference anyway? In: *ICAPS*. [S.l.: s.n.], 2009. p. 162–169. Citado na página [46](#).
- HOLTE, R. C. et al. Maximizing over multiple pattern databases speeds up heuristic search. *Artificial Intelligence*, Elsevier, v. 170, n. 16, p. 1123–1136, 2006. Citado 3 vezes nas páginas [11](#), [26](#), and [35](#).
- KNUTH, D. E. Estimating the efficiency of backtrack programs. *Mathematic of Computation*, v. 29, n. 129, p. 121–136, 1975. Citado na página [49](#).

KORF, R. E. Depth-first iterative-deepening: An optimal admissible tree search. *Artificial Intelligence*, v. 27, p. 97–109, 1985. Citado na página 31.

LELIS, L. H.; OTTEN, L.; DECHTER, R. Memory-efficient tree size prediction for depth-first search in graphical models. In: *Principles and Practice of Constraint Programming*. [S.l.]: Springer International Publishing, 2014. p. 481–496. Citado na página 59.

LELIS, L. H.; STERN, R.; STURTEVANT, N. R. Estimating search tree size with duplicate detection. In: *Seventh Annual Symposium on Combinatorial Search*. [S.l.: s.n.], 2014. Citado na página 52.

LELIS, L. H.; ZILLES, S.; HOLTE, R. C. Predicting the size of ida\* search tree. *Artificial Intelligence*, Elsevier, v. 196, p. 53–76, 2013. Citado 4 vezes nas páginas 31, 49, 50, and 59.

LELIS, L. H. Santana de. *Cluster-and-Conquer: A Paradigm for Solving State-Space Problems*. Tese (Doutorado) — University of Alberta, 2013. Citado na página 49.

NEMHAUSER, G. L.; WOLSEY, L. A.; FISHER, M. L. An analysis of approximations for maximizing submodular set functions—i. *Mathematical Programming*, Springer-Verlag, v. 14, n. 1, p. 265–294, 1978. Citado na página 44.

NISSIM, R.; HOFFMANN, J.; HELMERT, M. Computing perfect heuristics in polynomial time: On bisimulation and merge-and-shrink abstraction in optimal planning. In: *22nd International Joint Conference on Artificial Intelligence (IJCAI'11)*. [S.l.: s.n.], 2011. Citado na página 63.

PRIEDITIS, A. Machine discovery of effective admissible heuristics. *Machine Learning*, v. 12, p. 117–141, 1993. Disponível em: <<http://dx.doi.org/10.1007/BF00993063>>. Citado na página 49.

RAYNER, C.; STURTEVANT, N.; BOWLING, M. Subset selection of search heuristics. *Proceedings of the Twenty-Third International Joint Conference on Artificial Intelligence*, p. 637–643, 2013. Citado 5 vezes nas páginas 31, 43, 44, 62, and 69.

TOLPIN, D. et al. Towards rational deployment of multiple heuristics in a\*. In: AAAI PRESS. *Proceedings of the Twenty-Third international joint conference on Artificial Intelligence*. [S.l.], 2013. p. 674–680. Citado na página 26.

ZAHAVI, U. et al. Predicting the performance of ida\* using conditional distributions. *Journal of Artificial Intelligence Research*, v. 37, n. 1, p. 41–84, 2010. Citado na página 53.

# Index

sinopse de capítulo, [25](#), [31](#), [43](#), [59](#)